

ETC3250

Business Analytics

Week 1.
Introduction to Business Analytics & R

27 July 2015

Who are we?



Souhaib Ben Taieb



Di Cook



Rob Hyndman

Who are you?

- Aingaran Somaskanthan
- Andrew Ireland
- Braden Churcher
- Charles Condon
- Frances Kim
- Ivan Geerlings
- James Morrison
- Jin Zeng
- Joanna Pan
- Mitchell O'Hara-Wild
- Rohan Chandra
- Shin Wei Tan
- Stephanie Kobakian
- Vincent Tang
- ???

Outline

Week	Topic		Chapter	Lecturers
1	Introduction to business analytics & R	1		Rob, Souhaib
2	Statistical learning	2		Rob, Souhaib
3	Regression for prediction	3		Rob
4	Resampling	5		Rob
5	Dimension reduction	6,10		Rob, Souhaib
6	Visualization			Di
7	Visualization			Di
8	Classification	4,8		Souhaib, Di
9	Classification	4,9		Souhaib
10	Advanced classification	8		Di
11	Advanced regression	6		Di, Souhaib
12	Clustering	10		Souhaib, Di

Assessment

- Ten short weekly assignments, worth 2% each.
- One project due at the end of the semester, worth 20%.
- Exam in a computer lab (2 hours): 60%.

Task	Due Date	Value
Assignments 1–10	Wed 11:55pm each wk	2% each
Project	Fri 23 October	20%
Final exam	Official exam period	60%

Moodle site

- Includes all lecture notes, handouts, assignments
- Forum for asking questions, etc.
- No email please — use the forum.
- Assignment submissions

Key reference

**James, Witten, Hastie and Tibshirani
(2012) *An Introduction to Statistical Learning*. Springer.**

www.statlearning.com

- Free pdf online
- Data sets in associated R package **ISLR**
- R code for examples

What is business analytics?

Using **data** to gain **insights** and **understanding** of business problems and performance.

- Broader than **business intelligence** which focuses on describing and predicting performance.
- Broader than **econometrics** as we are interested in more than economics and finance.
- Narrower than **data science** as we are focusing on business issues.

What is business analytics?

Using **data** to gain **insights** and **understanding** of business problems and performance.

- Broader than **business intelligence** which focuses on describing and predicting performance.
- Broader than **econometrics** as we are interested in more than economics and finance.
- Narrower than **data science** as we are focusing on business issues.

What is business analytics?

“**Statistics** is the **science of learning from data**, and of **measuring, controlling, and communicating uncertainty**; [...].”

“Machine learning is a **scientific discipline** that explores the **construction and study of algorithms** that can **learn from data**”.

“Data mining, [...], is the **computational process** of **discovering patterns** in large data sets involving methods at the intersection of **artificial intelligence, machine learning, statistics, and database systems**”.

“Data Science means the **scientific study** of the **creation, validation and transformation of data** to **create meaning**”.

What is business analytics?

“Statistics is the **science of learning from data**, and of **measuring, controlling, and communicating uncertainty**; [...].”

“Machine learning is a **scientific discipline** that explores the **construction and study of algorithms** that can **learn from data**”.

“Data mining, [...], is the **computational process of discovering patterns** in large **data sets** involving methods at the intersection of **artificial intelligence, machine learning, statistics, and database systems**”.

“Data Science means the **scientific study** of the **creation, validation** and **transformation of data** to **create meaning**”.

What is business analytics?

“Statistics is the **science of learning from data**, and of **measuring, controlling, and communicating uncertainty**; [...].”

“Machine learning is a **scientific discipline** that explores the **construction and study of algorithms** that can **learn from data**”.

“Data mining, [...], is the **computational process** of **discovering patterns** in large **data sets** involving methods at the intersection of **artificial intelligence, machine learning, statistics, and database systems**”.

“Data Science means the **scientific study** of the **creation, validation** and **transformation of data** to **create meaning**”.

What is business analytics?

“Statistics is the **science of learning from data**, and of **measuring, controlling, and communicating uncertainty**; [...].”

“Machine learning is a **scientific discipline** that explores the **construction and study of algorithms** that can **learn from data**”.

“Data mining, [...], is the **computational process** of **discovering patterns** in large **data sets** involving methods at the intersection of **artificial intelligence, machine learning, statistics, and database systems**”.

“Data Science means the **scientific study** of the **creation, validation** and **transformation of data** to **create meaning**”.

What is business analytics?

“Statistics is the **science of learning from data**, and of **measuring, controlling, and communicating uncertainty**; [...].”

“Machine learning is a **scientific discipline** that explores the **construction and study of algorithms** that can **learn from data**”.

“Data mining, [...], is the **computational process** of **discovering patterns** in large **data sets** involving methods at the intersection of **artificial intelligence, machine learning, statistics, and database systems**”.

“Data Science means the **scientific study** of the **creation, validation** and **transformation of data** to **create meaning**”.

Why study business analytics?

false discoveries / wrong conclusions from data

Why study business analytics?

Some quotes

By 2018, the US could face a shortage of up to 190.000 workers with analytical skills —McKinsey

Top-ranked jobs: CareerCast

- 1 Actuary
- 2 Audiologist
- 3 Mathematician
- 4 Statistician
- 5 Biomedical Engineer
- 6 Data Scientist
- 7 Dental Hygienist
- 8 Software Engineer
- 9 Occupational Therapist
- 10 Computer Systems Analyst

Some quotes

By 2018, the US could face a shortage of up to 190.000 workers with analytical skills —McKinsey

Top-ranked jobs: CareerCast

- 1** Actuary
- 2** Audiologist
- 3** Mathematician
- 4** Statistician
- 5** Biomedical Engineer
- 6** Data Scientist
- 7** Dental Hygienist
- 8** Software Engineer
- 9** Occupational Therapist
- 10** Computer Systems Analyst

Some quotes

By 2018, the US could face a shortage of up to 190.000 workers with analytical skills —McKinsey

Top-ranked jobs: CareerCast

- 1** Actuary
- 2** Audiologist
- 3** Mathematician
- 4** Statistician
- 5** Biomedical Engineer
- 6** Data Scientist
- 7** Dental Hygienist
- 8** Software Engineer
- 9** Occupational Therapist
- 10** Computer Systems Analyst

Some quotes

Data Scientist: The Sexiest Job of the 21st Century —

Thomas H. Davenport and D. J. Patil, Harvard Business Review,
October 2012.

To have any hope of extracting anything useful from big data, . . . effective inferential skills are vital. That is, at the heart of extracting value from big data lies statistics — David J. Hand, 2014.

Most of my life I went to parties and heard a little groan when people heard what I did. Now they are all excited to meet me — Robert Tibshirani, a **statistics** professor at Stanford University, New York Times, January 26, 2012.

Am I a data scientist?

April 2013: Larry Wasserman blog

Data science: the end of statistics?

If you're analyzing data, you're doing statistics. You can call it data science or informatics or analytics or whatever, but it's still statistics.

normaldeviate.wordpress.com/2013/04/13/data-science-the-end-of-statistics/

Am I a data scientist?

April 2013: Larry Wasserman blog

Data science: the end of statistics?

If you're analyzing data, you're doing statistics. You can call it data science or informatics or analytics or whatever, but it's still statistics.

normaldeviate.wordpress.com/2013/04/13/data-science-the-end-of-statistics/

April 2013: Karl Broman blog

Data science is statistics

You may not like what some statisticians do. You may feel they don't share your values. They may embarrass you. But that shouldn't lead us to abandon the term "statistics".

kbroman.wordpress.com/2013/04/05/data-science-is-statistics/

Am I a data scientist?

April 2013: Larry Wasserman blog

Data science: the end of statistics?

If you're analyzing data, you're doing statistics. You can call it data science or informatics or analytics or whatever, but it's still statistics.

normaldeviate.wordpress.com/2013/04/13/data-science-the-end-of-statistics/

April 2013: Karl Broman blog

Data science is statistics

You may not like what some statisticians do. You may feel they don't share your values. They may embarrass you. But that shouldn't lead us to abandon the term "statistics".

kbroman.wordpress.com/2013/04/05/data-science-is-statistics/

July 2013: ASA President blog

Davidian: Aren't we data science?

Am I a data scientist?

November 2013: Andrew Gelman blog

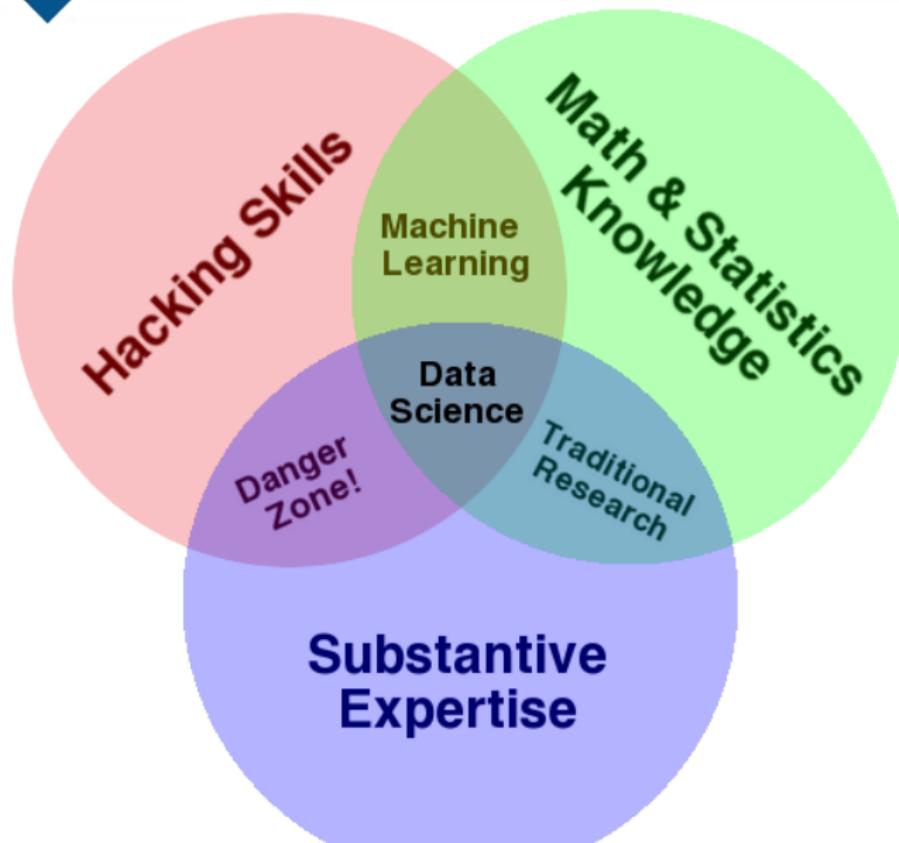
Statistics is the *least* important part of data science

There's so much that goes on with data that is about computing, not statistics. I do think it would be fair to consider statistics as a subset of data science ...

Statistics is important—don't get me wrong—statistics helps us correct biases ... estimate causal effects ... regularize so that we're not overwhelmed by noise ... fit models ... visualize data ... I love statistics! But it's not the most important part of data science, or even close.

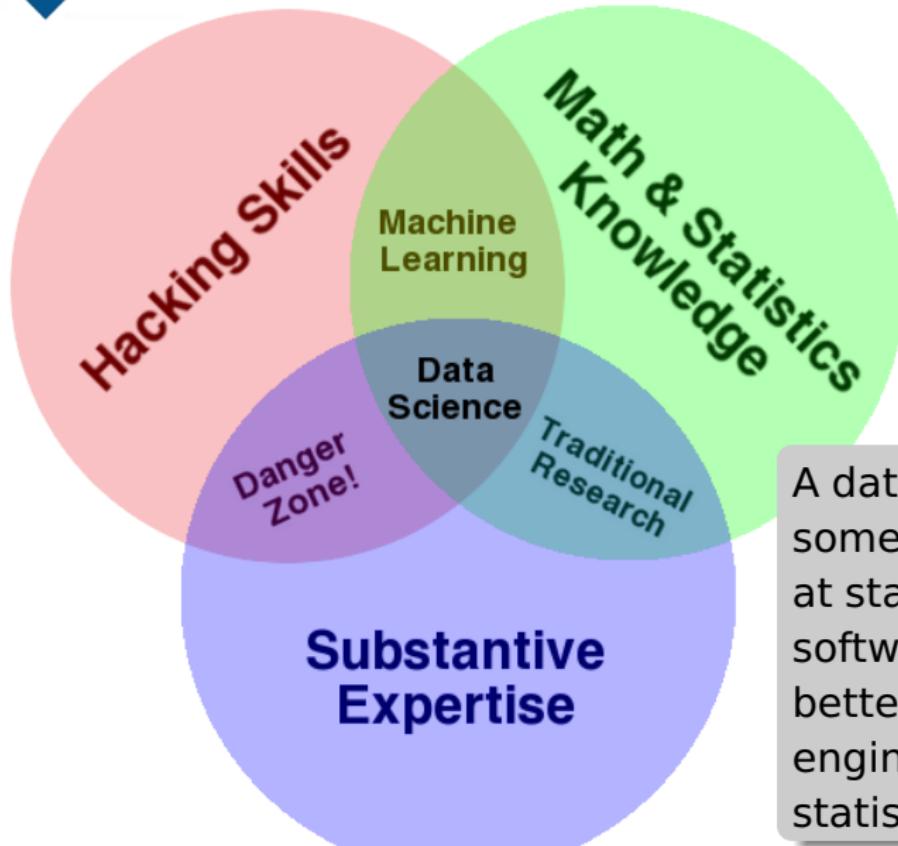
andrewgelman.com/2013/11/14/statistics-least-important-part-data-science/

Am I a data scientist?



Source: Drew Conway, Sept 2010. drewconway.com/zia/2013/3/26/the-data-science-venn-diagram. Reproduced under a Creative Commons License.

Am I a data scientist?

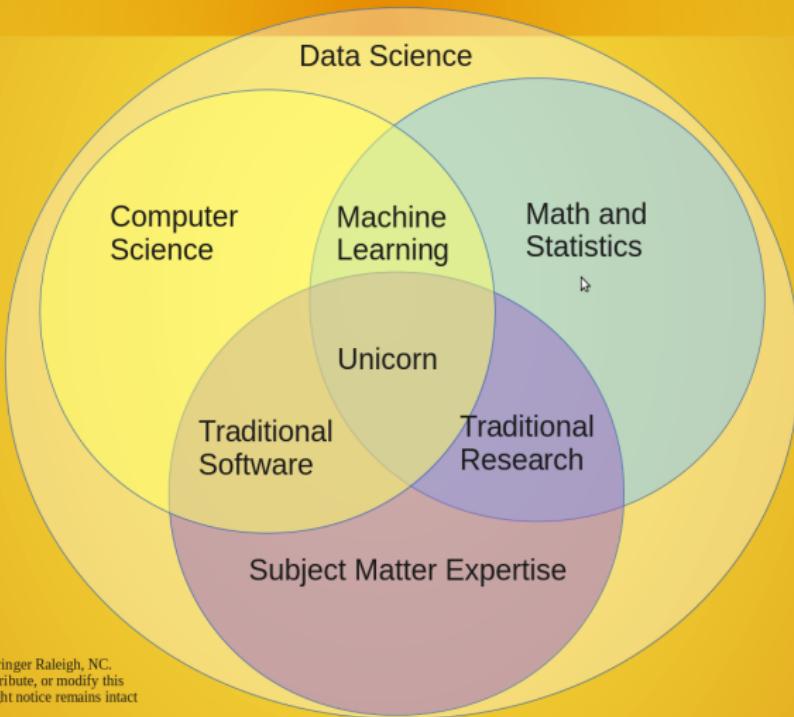


A data scientist is someone who is better at statistics than any software engineer and better at software engineering than any statistician.

Source: Drew Conway, Sept 2010. drewconway.com/zia/2013/3/26/the-data-science-venn-diagram. Reproduced under a Creative Commons License.

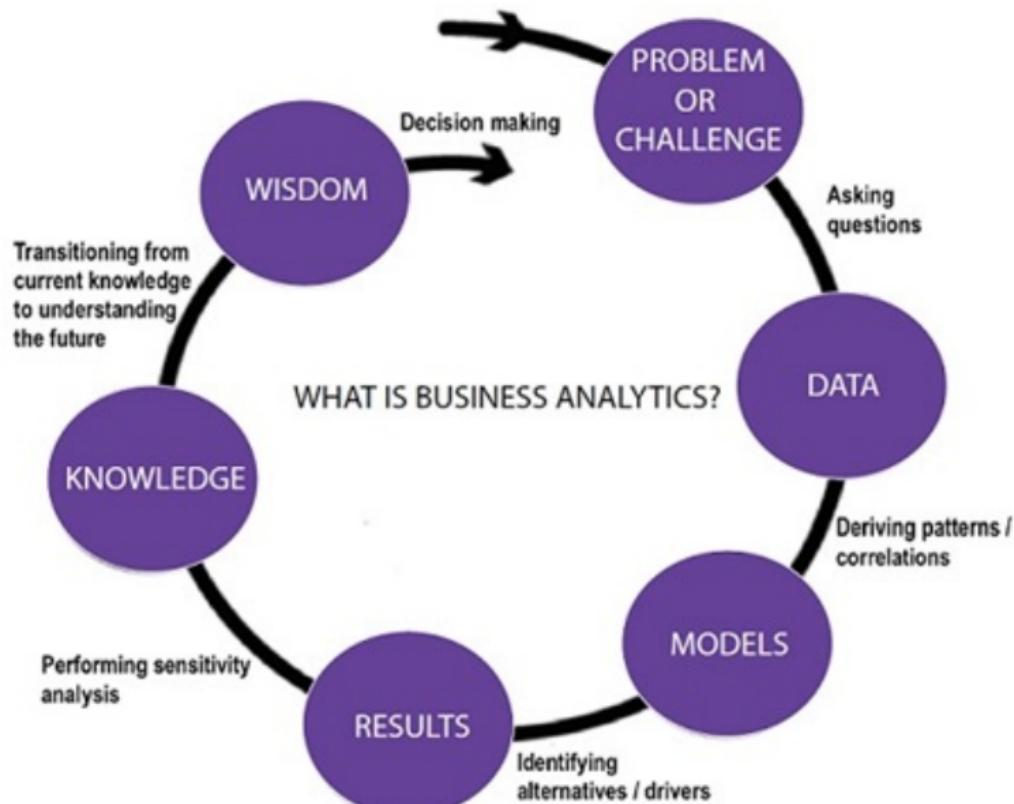
Am I a data scientist?

Data Science Venn Diagram v2.0



Copyright © 2014 by Steven Geringer Raleigh, NC.
Permission is granted to use, distribute, or modify this
image, provided that this copyright notice remains intact

The business analytics process



Source: <http://www.stern.nyu.edu/programs-admissions/executive-education/short-courses/schedule/short-course-program-7>

The business analytics tools

- Pulling together and cleaning data
- Exploring and visualizing data
- Fitting, comparing and assessing models
- Tools for fitting models: optimization, training and testing
- Tools for understanding randomness: simulation, resampling, permutation, cross-validation
- Tools for handling large data sets: dimension reduction, regularization, distributed computing.

Learning goals

- 1 Select and develop appropriate models for clustering, prediction or classification
- 2 Estimate and simulate from a variety of statistical models measure the uncertainty of a prediction or classification using resampling methods
- 3 Apply business analytic tools to produce innovative solutions in finance, marketing, economics and related areas
- 4 Manage very large data sets in a modern software environment explain and interpret the analyses undertaken clearly and effectively.

Teaching and learning approach

Two 1-hour lectures and a one 2 hour lab class each week for 12 weeks.

Learning goals

- 1 Select and develop appropriate models for clustering, prediction or classification
- 2 Estimate and simulate from a variety of statistical models measure the uncertainty of a prediction or classification using resampling methods
- 3 Apply business analytic tools to produce innovative solutions in finance, marketing, economics and related areas
- 4 Manage very large data sets in a modern software environment explain and interpret the analyses undertaken clearly and effectively.

Teaching and learning approach

Two 1-hour lectures and a one 2 hour lab class each week for 12 weeks.

R



R



RStudio

