



# ETC3250 Business Analytics: Data Wrangling

Souhaib Ben Taieb, Di Cook, Rob Hyndman

September 7, 2015

# Using the packages tidy, dplyr

During a ten week sensory experiment, 12 individuals were asked to assess taste of french fries (HOT CHIPS!) on several scales (how potato-y, buttery, grassy, rancid, paint-y do the fries taste?)

French fries were fried in one of three different oils, and each week individuals had to assess six batches of french fries (all three oils, replicated twice)


	##	time	treatment	subject	rep	potato	buttery	grassy	rancid
	## 61	1	1	3	1	2.9	0.0	0.0	0.0
	## 25	1	1	3	2	14.0	0.0	0.0	1.1
	## 62	1	1	10	1	11.0	6.4	0.0	0.0
	## 26	1	1	10	2	9.9	5.9	2.9	2.2
	## 63	1	1	15	1	1.2	0.1	0.0	1.1
	## 27	1	1	15	2	8.8	3.0	3.6	1.5

# This format is not ideal for data analysis

What code would be needed to plot each of the ratings over time as a different color?

```
library(ggplot2)
french_sub <- french_fries[french_fries$time == 10,]
qplot("1_potato", potato, data = french_sub,
      fill = I("red"), geom = "boxplot") +
  geom_boxplot(aes(x = "2_buttery", y = buttery),
    fill = I("orange")) +
  geom_boxplot(aes(x = "3_grassy", y = grassy),
    fill = I("yellow")) +
  geom_boxplot(aes(x = "4_rancid", y = rancid),
    fill = I("green")) +
  geom_boxplot(aes(x = "5_painty", y = painty),
    fill = I("blue")) +
  xlab("variable") + ylab("rating")
```

# The Plot



# What we have ..

We want to change this **wide format**:

--	--	--	--	--	--

# and what we want

to this **long format**:


- When gathering, you need to specify the **keys** (identifiers) and the **values** (measures).
- Keys/Identifiers: – Identify a record (must be unique) – Example: Indices on a random variable – Fixed by design of experiment (known in advance) – May be single or composite (may have one or more variables)
- Values/Measures: – Collected during the experiment (not known in advance) – Usually numeric quantities

# Gathering the French Fry Data

```
french_fries_long <- gather(french_fries, key = variable, value = rating)
```

```
head(french_fries_long)
```


```
##    time treatment subject rep variable rating
## 1     1           1       3    1  potato    2.9
## 2     1           1       3    2  potato   14.0
## 3     1           1      10    1  potato   11.0
## 4     1           1      10    2  potato    9.9
## 5     1           1      15    1  potato    1.2
## 6     1           1      15    2  potato    8.8
```



# Let's Re-write the code for our Plot

```
french_fries_long_sub <- french_fries_long[  
  french_fries_long$time == 10,]  
  
qplot(variable, rating, data = french_fries_long_sub,  
  fill = variable, geom = "boxplot")
```

# And plot it



# Long to Wide

In certain applications, we may wish to take a long dataset and convert it to a wide dataset (Perhaps displaying in a table).

##	time	treatment	subject	rep	variable	rating
## 1	1	1	3	1	potato	2.9
## 2	1	1	3	2	potato	14.0
## 3	1	1	10	1	potato	11.0
## 4	1	1	10	2	potato	9.9
## 5	1	1	15	1	potato	1.2
## 6	1	1	15	2	potato	8.8

# Spread

We use the **spread** function from tidyr to do this:

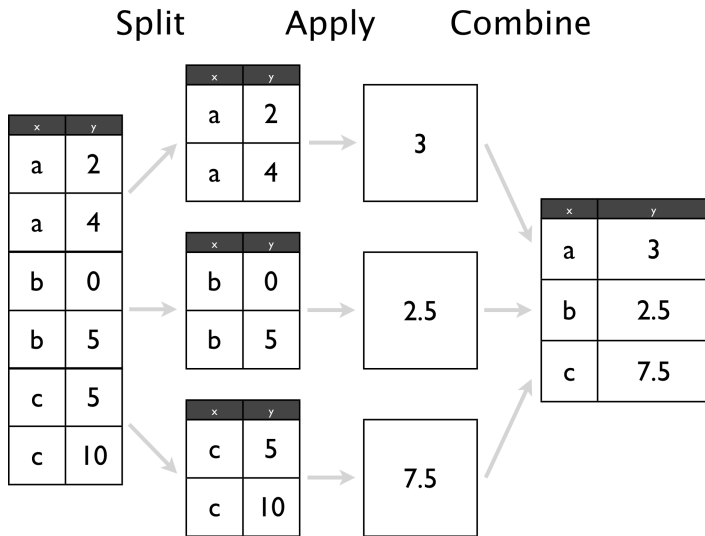
```
french_fries_wide <- spread(french_fries_long,  
  key = variable, value = rating)  
head(french_fries_wide)
```

```
##   time treatment subject rep potato buttery grassy rancid p  
## 1     1           1       3     1     2.9      0.0     0.0     0.0  
## 2     1           1       3     2    14.0      0.0     0.0     1.1  
## 3     1           1      10     1    11.0      6.4     0.0     0.0  
## 4     1           1      10     2     9.9      5.9     2.9     2.2  
## 5     1           1      15     1     1.2      0.1     0.0     1.1  
## 6     1           1      15     2     8.8      3.0     3.6     1.5
```

# The Split-Apply-Combine Approach

- *Split* a dataset into many smaller sub-datasets
- *Apply* some function to each sub-dataset to compute a result
- *Combine* the results of the function calls into a one dataset

# The Split-Apply-Combine Approach



# Split-Apply-Combine in dplyr

```
library(dplyr)
french_fries_split <- group_by(french_fries_long,
  variable) # SPLIT
french_fries_apply <- summarise(french_fries_split,
  m = mean(rating, na.rm = TRUE), s=sd(rating, na.rm=TRUE))
  # APPLY + COMBINE
french_fries_apply

## Source: local data frame [5 x 3]
##
##   variable      m      s
## 1  potato 6.9525180 3.584403
## 2  buttery 1.8236994 2.409758
## 3  grassy 0.6641727 1.320574
## 4  rancid 3.8522302 3.781815
## 5  painty 2.5217579 3.393717
```

# The pipe operator

- dplyr allows us to chain together these data analysis tasks using the `%>%` (pipe) operator
- `x %>% f(y)` is shorthand for `f(x, y)`
- Example:

```
french_fries %>%  
  gather(key = variable, value = rating, potato:painty) %>%  
  group_by(variable) %>%  
  summarise(rating = mean(rating, na.rm = TRUE))
```



There are five primary dplyr **verbs**, representing distinct data analysis tasks:

- Filter: Remove the rows of a data frame, producing subsets
- Arrange: Reorder the rows of a data frame
- Select: Select particular columns of a data frame
- Mutate: Add new columns that are functions of existing columns
- Summarise: Create collapsed summaries of a data frame

# Filter

```
french_fries %>%
```

```
  filter(subject == 3, time == 1)
```

```
##   time treatment subject rep potato buttery grassy rancid p
## 1     1         1       3   1    2.9     0.0    0.0    0.0
## 2     1         1       3   2   14.0     0.0    0.0    1.1
## 3     1         2       3   1   13.9     0.0    0.0    3.9
## 4     1         2       3   2   13.4     0.1    0.0    1.5
## 5     1         3       3   1   14.1     0.0    0.0    1.1
## 6     1         3       3   2    9.5     0.0    0.6    2.8
```

```
## Source: local data frame [5 x 2]
```

```
##
```

```
##   variable    rating
```

```
## 1  potato 6.9525180
```

```
## 2  buttery 1.8236994
```

```
## 3  grassy 0.6641727
```

```
## 4  rancid 2.8522202
```

# Arrange

```
french_fries %>%  
  arrange(desc(rancid)) %>%  
  head
```

##	time	treatment	subject	rep	potato	buttery	grassy	rancid	p
## 1	9	2	51	1	7.3	2.3	0	14.9	
## 2	10	1	86	2	0.7	0.0	0	14.3	
## 3	5	2	63	1	4.4	0.0	0	13.8	
## 4	9	2	63	1	1.8	0.0	0	13.7	
## 5	5	2	19	2	5.5	4.7	0	13.4	
## 6	4	3	63	1	5.6	0.0	0	13.3	

# Select

```
french_fries %>%  
  select(time, treatment, subject, rep, potato) %>%  
  head
```

```
##      time treatment subject rep potato  
## 61      1          1        3    1    2.9  
## 25      1          1        3    2   14.0  
## 62      1          1       10    1   11.0  
## 26      1          1       10    2    9.9  
## 63      1          1       15    1    1.2  
## 27      1          1       15    2    8.8
```

# Summarise

```
french_fries %>%  
  group_by(time, treatment) %>%  
  summarise(mean_rancid = mean(rancid), sd_rancid = sd(rancid))
```

```
## Source: local data frame [30 x 4]
```

```
## Groups: time
```

```
##
```

##	time	treatment	mean_rancid	sd_rancid
## 1	1	1	2.758333	3.212870
## 2	1	2	1.716667	2.714801
## 3	1	3	2.600000	3.202037
## 4	2	1	3.900000	4.374730
## 5	2	2	2.141667	3.117540
## 6	2	3	2.495833	3.378767
## 7	3	1	4.650000	3.933358
## 8	3	2	2.895833	3.773532
## 9	3	3	3.600000	3.592867

# Dates and Times

- Dates are deceptively hard to work with in R.

**Example:** 02/05/2012. Is it February 5th, or May 2nd?

Other things are difficult too:

- Time zones
- POSIXct format in base R is challenging

The **lubridate** package helps tackle some of these issues.

# Basic Lubridate Use

```
library(lubridate)
```

```
now()
```

```
today()
```

```
now() + hours(4)
```

```
today() - days(2)
```

```
## [1] "2015-09-07 06:50:40 AEST"
```

```
## [1] "2015-09-07"
```

```
## [1] "2015-09-07 10:50:40 AEST"
```

```
## [1] "2015-09-05"
```

# Parsing Dates

```
ymd("2013-05-14")
mdy("05/14/2013")
dmy("14052013")
ymd_hms("2015:05:14 14:50:30", tz = "America/Chicago")
ymd_hms("2015:05:14 14:50:30", tz = "Australia/Melbourne")
today(tzone = "America/Chicago")
today(tzone = "Australia/Melbourne")
```

```
## [1] "2013-05-14 UTC"
## [1] "2013-05-14 UTC"
## [1] "2013-05-14 UTC"
## [1] "2015-05-14 14:50:30 CDT"
## [1] "2015-05-14 14:50:30 AEST"
## [1] "2015-09-06"
## [1] "2015-09-07"
```



# Dates example: Oscars date of birth

```
oscars <- read.csv("../data/oscars.csv", stringsAsFactors=FALSE)
summary(oscars$DOB)
head(oscars$DOB)
oscars$DOB <- as.Date(oscars$DOB, format="%m/%d/%Y")
summary(oscars$DOB)
```

```
##      Length      Class      Mode 
##      423 character character
```

```
## [1] "9/30/1895" "7/23/1884" "4/23/1894" "10/6/2006" "2/2/1884"
```

```
##           Min.        1st Qu.          Median            Mean            3rd Qu. 
## "1868-04-10" "1934-09-18" "1957-06-23" "1962-05-21" "2008-05-21" 
##           Max. 
## "2029-12-13"
```

# Calculating on dates

- You should never ask a woman her age, but ... really!

```
oscars$DOByr <- year(oscars$DOB)
summary(oscars$DOByr)
oscars %>% filter(DOByr == "2029") %>% select(Name, Sex, DOB)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1868	1934	1957	1962	2008	2029

##	Name	Sex	DOB
## 1	Audrey Hepburn	Female	2029-05-04
## 2	Grace Kelly	Female	2029-11-12
## 3	Miyoshi Umeki	Female	2029-04-03
## 4	Christopher Plummer	Male	2029-12-13


# Months

```
oscars$DOBmth <- month(oscars$DOB, )  
table(oscars$DOBmth)  
oscars$DOBmth <- factor(oscars$DOBmth, levels=1:12,  
  labels=month.abb)
```

```
##  
##  1  2  3  4  5  6  7  8  9 10 11 12  
## 32 35 37 53 35 27 37 33 31 31 30 42
```


# Now plot it

```
qplot(DOBmth, data=oscars)
```



# Should you be born in April?

```
df <- data.frame(m=sample(1:12, 423, replace=TRUE))  
df$m2 <- factor(df$m, levels=1:12,  
  labels=month.abb)  
qplot(m2, data=df)
```



# Working with strings

- Example: NBA salaries
- ESPN provides basketball players' salaries for the 2013-2014 season at <http://espn.go.com/nba/salaries>

```
##      RK      NAME      TEAM      SALARY
## 1  1      Kobe Bryant, SG  Los Angeles Lakers $25,000,000
## 2  2      Joe Johnson, SF   Brooklyn Nets $24,894,863
## 3  3      LeBron James, SF  Cleveland Cavaliers $22,970,500
## 4  4      Carmelo Anthony, SF New York Knicks $22,875,000
## 5  5      Dwight Howard, C   Houston Rockets $22,359,364
## 6  6      Chris Bosh, C      Miami Heat $22,192,730

## 'data.frame':    423 obs. of  4 variables:
##  $ RK      : Factor w/ 394 levels "1","10","11",...: 1 12 23 3
##  $ NAME     : Factor w/ 394 levels "Al Jefferson, C",...: 31 24
##  $ TEAM     : Factor w/ 32 levels "Atlanta Hawks",...: 14 3 6 1
##  $ SALARY   : Factor w/ 299 levels "$13,400,000",...: 35 34 33
```

# Cleaning NBA salaries data

```
head(nba14$SALARY)
```

```
# get rid of $ and , in salaries and convert to numeric:
```

```
gsub("$,", "", head(as.character(nba14$SALARY)))
```

```
nba14$SALARY <- as.numeric(gsub("$,", "",  
  as.character(nba14$SALARY)))
```

```
## [1] $25,000,000 $24,894,863 $22,970,500 $22,875,000 $22,359,364
```

```
## 299 Levels: $13,400,000 $13,437,500 $13,500,000 $13,600,000
```

```
## [1] "25000000" "24894863" "22970500" "22875000" "22359364"
```

```
## Warning: NAs introduced by coercion
```

- Where does the warning come from?

# Cleaning NBA salaries data: hunting the warning

```
head(subset(nba14, is.na(SALARY)))
```

```
##      RK NAME TEAM SALARY
## 11 RK NAME TEAM      NA
## 22 RK NAME TEAM      NA
## 33 RK NAME TEAM      NA
## 54 RK NAME TEAM      NA
## 65 RK NAME TEAM      NA
## 76 RK NAME TEAM      NA
```

- We don't need these rows - delete all of them

```
dim(nba14)
nba14 <- nba14[-which(nba14$RK=="RK"),]
dim(nba14)
```

```
## [1] 423    4
```



# Cleaning NBA data

- Separate names into first, last, and position

```
library(stringr)
splits <- str_split(as.character(nba14$NAME), pattern="(", )| '
splits[1:3]
library(plyr)
numnames <- ldply(splits, length)
summary(numnames) # some players have multiple names, ... sigh

## [[1]]
## [1] "Kobe"      "Bryant" "SG"
##
## [[2]]
## [1] "Joe"       "Johnson" "SF"
##
## [[3]]
## [1] "LeBron" "James" "SF"
```

# Cleaning data

- There's only limited possibilities in terms of what we can do automatically about people with multiple names - we will deal with them alongside the other ones and flag them ... maybe we should leave first and last name together.

```
head(splits[numnames>3], 5)
```

```
sum(numnames>3)
```

```
## [[1]]
```

```
## [1] "Otto"      "Porter" "Jr."      "SF"
```

```
##
```

```
## [[2]]
```

```
## [1] "Frank"      "Kaminsky" "III"      "C"
```

```
##
```

```
## [[3]]
```

```
## [1] "Kelly" "Oubre" "Jr."      "SF"
```

```
##
```

```
## [[4]]
```

# Cleaning NBA data

```
splitnames <- ldply(splits, function(x)
  c(name=paste(x[-length(x)], collapse=" "),
    position=x[length(x)]))
head(splitnames) # looks OK
# now copy into the nba14 data frame
nba14$name <- as.factor(splitnames$name)
nba14$position <- as.factor(splitnames$position)
```

```
##           name position
## 1      Kobe Bryant      SG
## 2      Joe Johnson      SF
## 3    LeBron James      SF
## 4 Carmelo Anthony      SF
## 5   Dwight Howard       C
## 6      Chris Bosh       C
```

# Cleaned data ...?


```
summary(nba14)
```

```
##           RK           NAME
##  1      :  1  Al Jefferson, C      :  1  Portland Trail BL
## 10      :  1  Andrew Bogut, C      :  1  Boston Celtics
## 11      :  1  Blake Griffin, PF    :  1  Brooklyn Nets
## 12      :  1  Carmelo Anthony, SF  :  1  Charlotte Hornets
## 13      :  1  Chandler Parsons, SF :  1  Toronto Raptors
## 14      :  1  Chris Bosh, C        :  1  Atlanta Hawks
## (Other):387  (Other)              :387  (Other)
##           SALARY           name      position
## Min.      : 525093  Aaron Brooks  :  1  C :57
## 1st Qu.: 1499187  Aaron Gordon  :  1  PF:90
## Median : 3283181  Aaron Harrison:  1  PG:74
## Mean      : 5267665  Adreian Payne :  1  SF:82
## 3rd Qu.: 7000000  Al Horford    :  1  SG:90
## Max.      :25000000  Al Jefferson  :  1
```

# Cleaned data ...?

- Numbers might still be wrong, but now we are in a position to check for that.

```
library(ggplot2)
qplot(SALARY, geom="histogram", data=nba14)
```



# Show it

```
qplot(position, SALARY, geom="boxplot", data=nba14)
```

